



# **CLOUD COMPUTING APPLICATIONS**

Serverless Cloud Computing:  
Landscape

Prof. Reza Farivar

# Serverless Cloud Computing Landscape

- Different Categories of Serverless Computing
  - Compute
    - Platform as a Service (Apps)
    - Function as a Service (Functions)
    - Container as a Service (Containers)
  - Storage
    - Blobs (Binary Large Objects)
    - Key/Value Datastores
  - Analytics
  - AI and ML

# Compute: PaaS

- The unit of compute is a full App
- You still select which server / platform you need, only that you don't need to manage it anymore
- Amazon Elastic BeanStalk
- Google AppEngine
- Microsoft Azure App Service
- IBM CloudFoundry
  - Based on Open Source CloudFoundry
    - Originally developed by VMware, transferred to Pivotal Software (a joint venture by EMC, VMware and General Electric), brought back into VMware at the end of 2019
- Oracle Java Cloud Service

# Function as a Service

- This is what most people think of as “Serverless”
- Unit of compute is a function
  - Functions running when “events” are triggered
- Amazon AWS Lambda
- Microsoft Azure Functions
- Google Cloud Functions
- IBM Cloud Functions
  - Based on Apache Open Whisk
- Oracle Functions
  - Apache Fn
- Open Source Open Lambda

# Container as a Service

- Function as a Service on HEAVY steroids!
  - Containers have to be Stateless
  - Unit of compute is a whole container
  - Container running when “events” are triggered
- Amazon Elastic Container Service (ECS), Elastic Kubernetes Service (EKS), Fargate
- Microsoft Azure Kubernetes Service (AKS)
  - Built on top of Open Source KEDA
- Google Cloud Run, Anthos
  - Built on top of Open Source Knative
    - Kubernetes-based platform to deploy and manage modern serverless workloads.
- IBM Cloud Kubernetes Service
- Oracle Container Engine for Kubernetes
- Open Source Kubernetes
  - Google Borg

# Serverless Storage: Blobs

- Blob Storage
  - Amazon Simple Storage Service (S3)
  - Microsoft Azure Blob Storage
  - IBM Object Storage
  - Google Cloud Storage
  - Oracle Object Storage

# Serverless Storage: Key/Value Database

- Distributed NOSQL key/value storage service
- Amazon DynamoDB
- Microsoft Azure Cosmos DB
- Google Cloud Datastore, Firestore, Cloud BigTable
- IBM Cloudant
- Open Source: Cassandra



# **CLOUD COMPUTING APPLICATIONS**

## Serverless Architecture


Roy Campbell & Reza Farivar



# Introduction to Serverless Architecture

- “Applications where some amount of server-side logic is still written by the application developer but unlike traditional architectures is run in stateless compute containers that are event-triggered, ephemeral (may only last for one invocation), and fully managed by a 3rd party”
- ‘Functions as a service / FaaS’
- AWS Lambda is one of the most popular implementations of FaaS at present, but there are others

# Introduction to Serverless Architecture



# Desktop Platform


Microsoft

Operating System

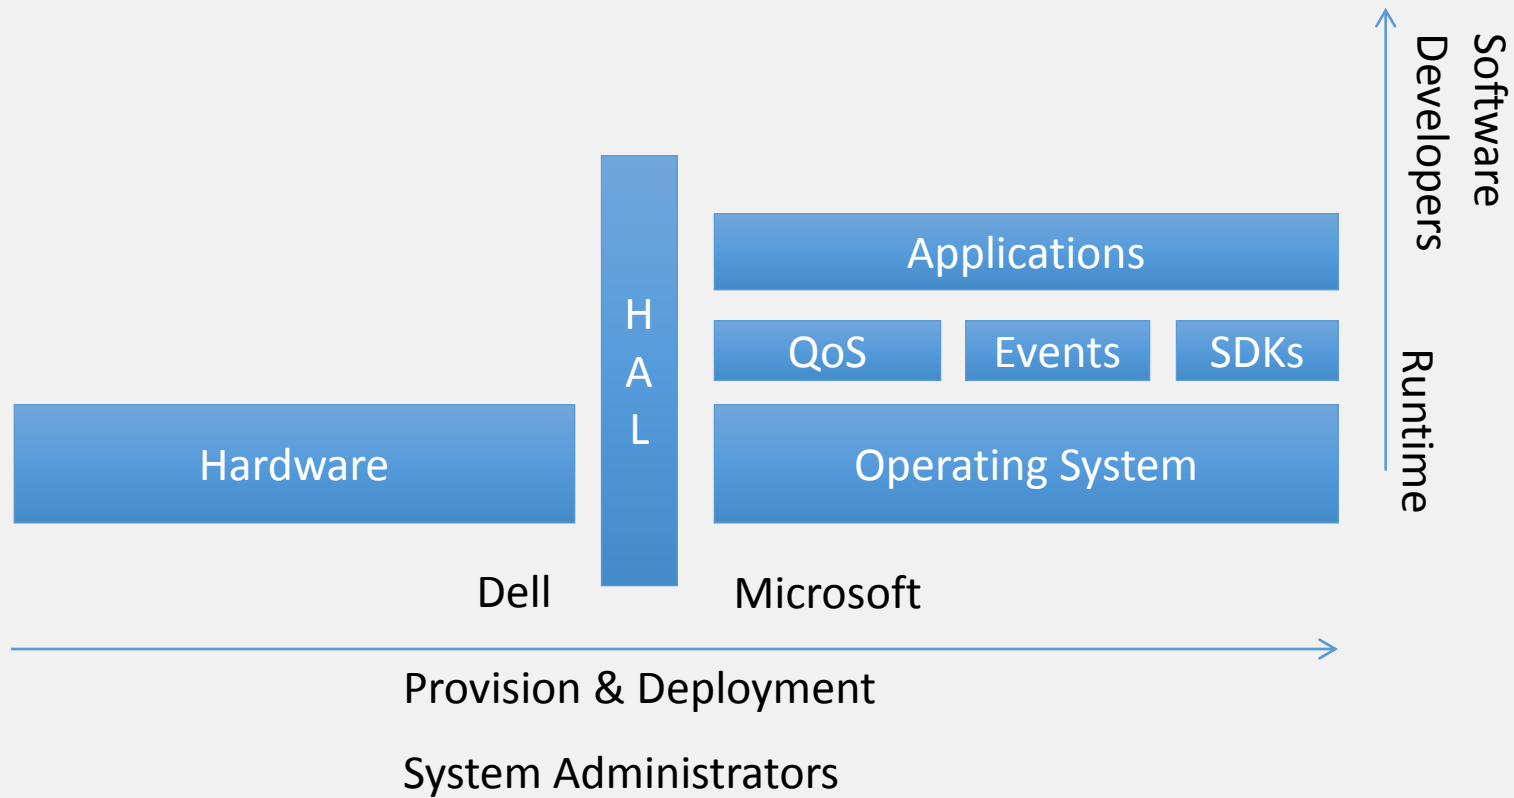
HAL

Dell

Hardware




# Desktop Platform




# AWS Cloud Platform, 2010

AWS (IaaS)



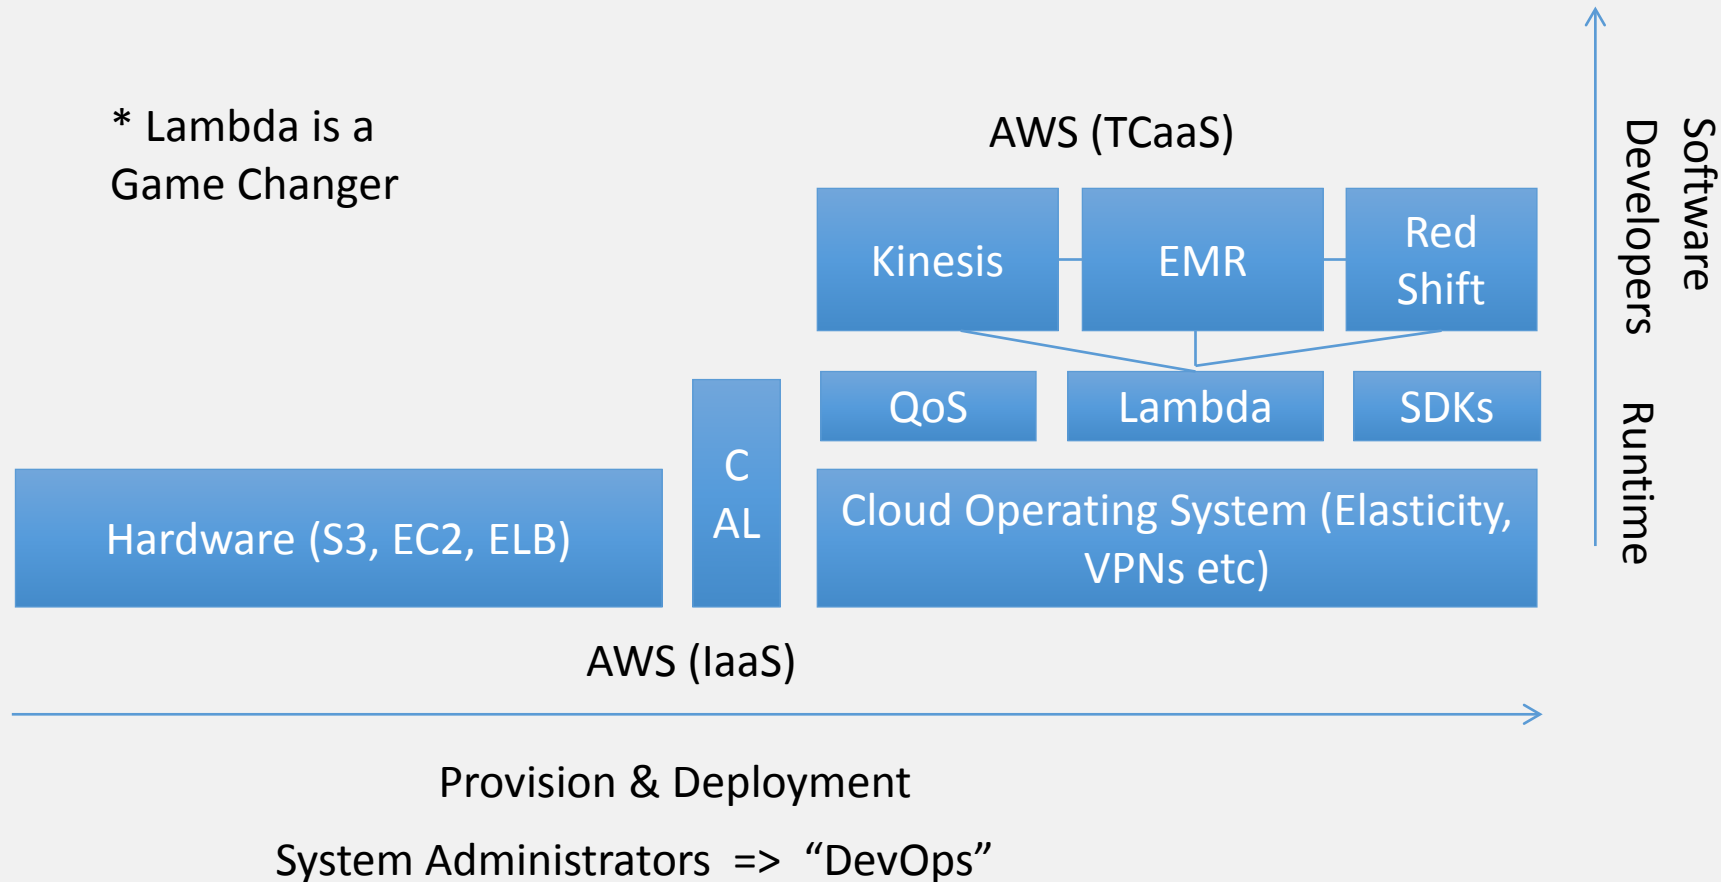
Provision & Deployment  
System Administrators

# AWS Cloud platform, 2014



# AWS Cloud Platform 2016

\* Lambda is a  
Game Changer



# AWS Elastic BeanStalk

- Deploy and scale web applications easily
- Languages: Java, .NET, PHP, Node.js, Python, Ruby, Docker
- Servers: Apache, Nginx, Phusion Passenger, IIS
- Simply upload your code; AWS handles:

Deployment

Auto scaling

Capacity Provisioning

Health Monitoring

Load balancing



# AWS Lambda Event-driven Compute

- Runs stateless, request-driven code called **Lambda functions in Java, NodeJS & Python**
- Triggered by events (state transitions) in other AWS services
- Pay only for the requests served and the compute time
- Focus on business logic, not infrastructure.
- Just upload your code; AWS Lambda handles:

Capacity

Monitoring

Fault Tolerance

Scaling

Logging

Security Patching

Deployment

Web service front end

# AWS Lambda Event Sources

- Amazon S3
- Amazon DynamoDB
- Amazon Kinesis Streams
- Amazon Simple Notification Service
- Amazon Simple Email Service
- Amazon Cognito
- AWS CloudFormation
- Amazon CloudWatch Logs
- Amazon CloudWatch Events
- Scheduled Events (powered by Amazon CloudWatch Events)
- AWS Config
- Amazon Echo
- Amazon API Gateway
- Other Event Sources: Invoking a Lambda Function On Demand
- Sample Events Published by Event Sources

# AWS Lambda Execution Environment

- State-less functions
- You can use multi-threading, etc.
- 500 MB of /tmp storage space
- You set how much memory you need:
  - From 128 MB to 1.5GB
  - 64GB increments
  - CPU scales accordingly
- Function should finish in a certain time
  - Default 3 seconds, up to 300 seconds

# AWS Lambda Pricing

- You pay per use of your function
- \$0.20 per 1 million function call
- Also, \$0.00001667 for every GB-second used



# **CLOUD COMPUTING APPLICATIONS**

Amazon S3 BLOB Storage

Roy Campbell & Reza Farivar

# Definition

Online file storage web service offered by Amazon Web Services. Amazon S3 provides storage through web service interfaces REST, SOAP, BitTorrent. (wikipedia)

<https://aws.amazon.com/s3/>

# Use case

- Scalability, high availability, low latency – 99.99% availability
  - Files up to 5 terabytes
  - Objects stored in buckets owned by users
  - User assigned keys refer to objects
- 
- Amazon Machine Images (exported as a bundle of objects)
  - SmugMug, Hadoop file store, Netflix, reddit, Dropbox, Tumbler

# Simple Storage Service (S3)

- A **bucket** is a container for objects and describes location, logging, accounting, and access control.
  - A bucket has a name that must be **globally unique**.
  - `http://bucket.s3.amazonaws.com`
  - `http://bucket.s3-aws-region.amazonaws.com`.
- A bucket can hold any number of **objects**, which are files of up to 5TB.
  - `http://bucket.s3.amazonaws.com/object`
  - `http://johnsmith.s3.amazonaws.com/photos/puppy.jpg`



# Fundamental operations corresponding to HTTP actions:

`http://bucket.s3.amazonaws.com/object`

- POST a new object or update an existing object.
- GET an existing object from a bucket.
- DELETE an object from the bucket
- LIST keys present in a bucket, with a filter.

A bucket has a **flat directory structure**

# S3 Weak Consistency Model

“Updates to a single key are **atomic**....”

“Amazon S3 achieves high availability by replicating data across multiple servers within Amazon's data centers. If a PUT request is successful, your data is safely stored. However:

- A process writes a new object to Amazon S3 and immediately attempts to read it. Until the change is fully propagated, Amazon S3 might report "key does not exist."
- A process writes a new object to Amazon S3 and immediately lists keys within its bucket. Until the change is fully propagated, the object might not appear in the list.
- A process replaces an existing object and immediately attempts to read it. Until the change is fully propagated, Amazon S3 might return the prior data.
- A process deletes an existing object and immediately attempts to read it. Until the deletion is fully propagated, Amazon S3 might return the deleted data.”

# S3 Command Line Interface

```
aws s3 mb s3://bucket
...    cp localfile s3://bucket/key
      mv s3://bucket/key s3://bucket/newname
      ls s3://bucket
      rm s3://bucket/key
      rb s3://bucket

aws s3 help
aws s3 ls help
```




# **CLOUD COMPUTING APPLICATIONS**

Serverless Storage: DynamoDB

Prof. Reza Farivar

# DynamoDB

- DynamoDB is a fully managed NoSQL database provided by Amazon AWS
- Think of it as a massive distributed B-Tree data structure in the cloud
  - Accessing specific items is blazingly fast
- Distributed system
  - Using the consistent hashing algorithm in a ring



# Usage model

- First create a table
  - Remember that it's a managed service, so just create a table using the console (or CLI, API)
    - We will use the Python Boto3 package in this lesson
- While creating the table, define the primary key
  - This key will be used by DynamoDB to distribute key/values in different partitions
- Optionally, identify a sort key
  - The sort key is used to keep the items in a partition sorted
  - Will be useful for query and scan later on



# Using the table: Put

- Having defined the table, we can now put values into it.
  - Note: DynamoDB items are limited to 400KB size

```
import boto3
dynamodb = boto3.resource('dynamodb')
table = dynamodb.Table('users')

table.put_item(
    Item={
        'username': 'janedoe',
        'first_name': 'Jane',
        'last_name': 'Doe',
        'age': 25,
        'account_type': 'standard_user',
    }
)
```

# Using the table: Get

- Retrieve an item

```
response = table.get_item(  
    Key={  
        'username': 'janedoe',  
        'last_name': 'Doe'  
    }  
)  
item = response['Item']  
print(item)
```

#Expected output:

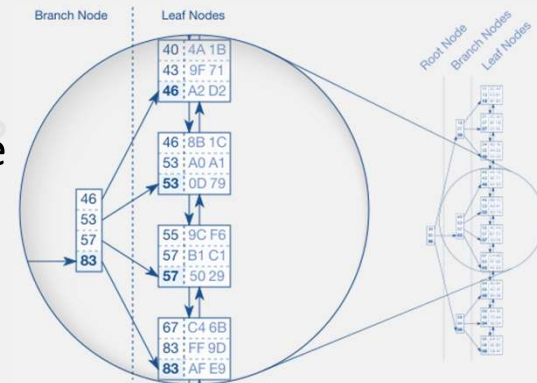
```
{u'username': u'janedoe',  
 u'first_name': u'Jane',  
 u'last_name': u'Doe',  
 u'account_type': u'standard_user',  
 u'age': Decimal('25')}
```



# Query and Scanning

- In RDBMS world\*, query is usually defined as an operation where there is a usable index available, and we can quickly retrieve the item in  $\text{Log}(n)$  time
- In comparison, scan happens where there is no usable index, and the engine has to read every record and test for a condition
- An RDBMS engine parses a SELECT statement, and performs query optimization, all behind the curtain
- DynamoDB just allows you to be your own database engine!

\*See <https://use-the-index-luke.com/sql/where-clause/searching-for-ranges/greater-less-between-tuning-sql-access-filter-predicates>



# Query

- Query only works on the primary key already defined for the table
- Or any other attribute that we have explicitly made a secondary index for it
- If a composite primary key was used (hash key + sort key), we can ask query to return a conditional range of value

```
response = table.query(  
    KeyConditionExpression=Key('username').eq('johndoe')  
)  
items = response['Items']  
print(items)
```

#Expected output:

```
[{'username': u'johndoe',  
  u'first_name': u'John',  
  u'last_name': u'Doe',  
  u'account_type': u'standard_user',  
  u'age': Decimal('25'),  
  u'address': {'city': u'Los Angeles',  
               u'state': u'CA',  
               u'zipcode': Decimal('90001'),  
               u'road': u'1 Jefferson Street'}}]
```

# Scan

- What if we want to perform a query conditioned on attributes that there is no index for them?
- Scan will return everything!
  - It allows to filter based on any arbitrary condition

```
response = table.scan(  
    FilterExpression=Attr('age').lt(27)  
)  
items = response['Items']  
print(items)
```

# Secondary Index

- Similar to the main index, requires a partition key and a sort key
- Local (LSI)
  - First released by Amazon in 2013
  - Immediately consistent
  - Once created, the table cannot grow any more
    - All the records that share the same partition key need to fit in 10GB
    - Once the allocation is full, writes fail ☹️
- Global (GSI)
  - Came a few months after local indexes
    - Eventual consistency model
    - Do not constrain table size 😊



# **CLOUD COMPUTING APPLICATIONS**

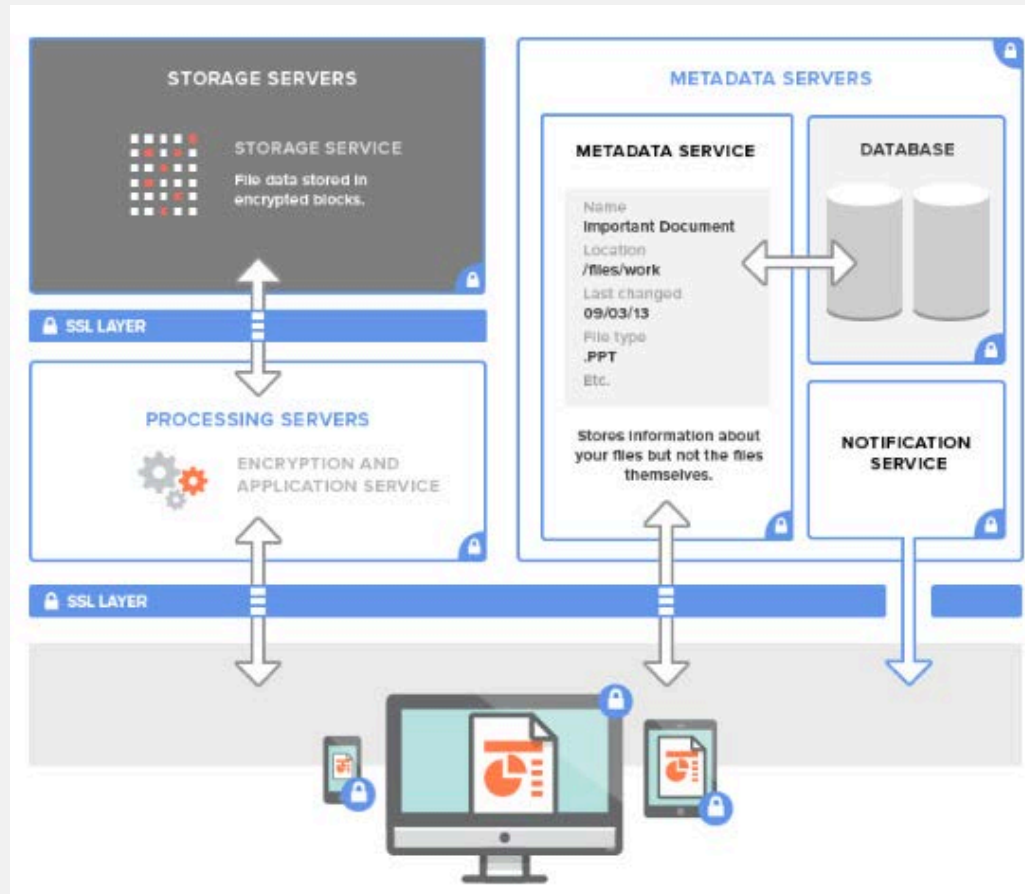
Dropbox Cloud API

Roy Campbell & Reza Farivar

# Cloud Storage

- One interesting case study is Dropbox
- Dropbox offers cloud file storage
  - Easily synced across multiple devices
  - Accessible through web interface, mobile apps, and directly integrated with the file system on PCs
- Dropbox itself uses clouds!
  - Metadata stored in Dropbox servers
  - Actual files stored in Amazon S3
  - Amazon EC2 instances run the logic

# Dropbox Architecture



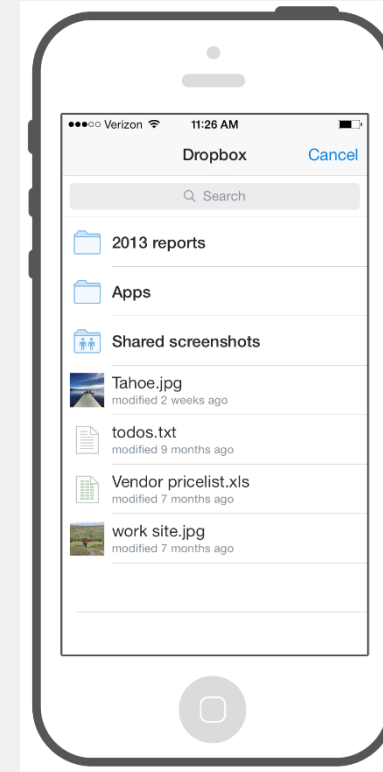

# Dropbox API

- Two levels of API access to Dropbox
  - Drop-ins
    - Cross-platform UI components that can be integrated in minutes
    - *Chooser* allows instant access to files in Dropbox
    - *Saver* makes saving files to Dropbox easy
  - Core API
    - Support for advanced functionality like search, revisions, and restoring file
    - Better fit for deeper integration



# Drop-In API

- Simple objects
  - *Chooser* available for JavaScript, Android and iOS
  - *Saver* on web and mobile web
- Handles all the authentication (OAuth), file browsing
- Chooser object returns the following:
  - Link: URL to access the file
  - File name
  - File Size
  - Icon
  - Thumbnails
- Saver
  - Pass in URL, filename and options



# Core API

- Many languages and environments
  - Python, Ruby, PHP, Java, Android, iOS, OS X, HTTP
- Based on HTTP and OAuth
  - OAuth v1, OAuth v2
- Low-level calls to access and manipulate a user's Dropbox account
  - Create URL schemes
  - Upload files
  - Download files
  - List files and folders
  - Delta
  - Metadata access
  - Create and manage file sharing